

# cf-01-rzepinski-ratings\_analysis

January 21, 2019

## 1 Important

make data has to be run before running any notebook cell

## 2 Imports

```
In [1]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import matplotlib
import numpy as np
```

```
In [2]: book_df = pd.read_csv('../data/raw/book.csv')
ratings_df = pd.read_csv('../data/raw/ratings.csv')
```

```
In [3]: ratings_df.head(1)
```

```
Out[3]:
```

	user_id	book_id	rating
0	1	258	5

Make sure there are no duplicates in ratings.

```
In [4]: ratings_df[ratings_df.duplicated(['user_id', 'book_id'], keep=False)]
```

```
Out[4]: Empty DataFrame
Columns: [user_id, book_id, rating]
Index: []
```

## 3 Visualization settings

```
In [5]: sns.set(context='paper', font_scale=1.2, style='ticks', palette='muted',
rc={"axes.labelsize":16, "ytick.labelsize": 14, "xtick.labelsize":14,
"font.family": "sans-serif"})
```

## 4 Ratings user and book coverage

```
In [6]: ratings_df.groupby('user_id')['book_id'].count().describe()
```

```
Out [6]: count      53424.000000
         mean       111.868804
         std        26.071224
         min        19.000000
         25%        96.000000
         50%       111.000000
         75%       128.000000
         max       200.000000
         Name: book_id, dtype: float64
```

All users rated at least 19 books. Such situation is rarely encountered in similar datasets.

```
In [7]: ratings_df.groupby('book_id')['user_id'].count().describe()
```

```
Out [7]: count      10000.000000
         mean       597.647900
         std       1267.289788
         min        8.000000
         25%       155.000000
         50%       248.000000
         75%       503.000000
         max      22806.000000
         Name: user_id, dtype: float64
```

All books have been rated at least 8 times.


## 5 How users rate books?

```
In [8]: ratings_df['rating'].describe()
```

```
Out [8]: count      5.976479e+06
         mean      3.919866e+00
         std       9.910868e-01
         min       1.000000e+00
         25%       3.000000e+00
         50%       4.000000e+00
         75%       5.000000e+00
         max       5.000000e+00
         Name: rating, dtype: float64
```

```
In [9]: sns.countplot(ratings_df.rating)
```

```
Out [9]: <matplotlib.axes._subplots.AxesSubplot at 0x7f04c22dd320>
```



```
In [10]: ratings_df.groupby('user_id')['rating'].mean().describe()
```

```
Out[10]: count      53424.000000
         mean         3.928512
         std         0.449543
         min         1.000000
         25%         3.633929
         50%         3.920455
         75%         4.223214
         max         5.000000
         Name: rating, dtype: float64
```

```
In [11]: len(ratings_df.groupby('user_id').filter(lambda x: x['rating'].mean() ==
         0.0)['user_id'].unique())
```

```
Out[11]: 0
```

```
In [12]: len(ratings_df.groupby('user_id').filter(lambda x: x['rating'].mean() ==
         5.0)['user_id'].unique())
```

```
Out[12]: 266
```

```
In [13]: user_mean_ratings_plot = sns.distplot(ratings_df.groupby('user_id')['rating'].mean(),
         kde=False)
         user_mean_ratings_plot.set(xlabel='Ratings mean', ylabel='Frequency')
```


```
/home/rzepinski/Documents/Inzynierka/Recommendation-system/rs-
venv/lib/python3.7/site-packages/scipy/stats/stats.py:1713: FutureWarning: Using a
non-tuple sequence for multidimensional indexing is deprecated; use `arr[tuple(seq)]`
```

```

instead of `arr[seq]`. In the future this will be interpreted as an array index,
`arr[np.array(seq)]`, which will result either in an error or a different result.
    return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval

```

Out[13]: [Text(0, 0.5, 'Frequency'), Text(0.5, 0, 'Ratings mean')]



People rate differently - some give only 5 stars reviews, some are more harsh than others, for some only perfect book should get 5 star rating and so on. Generally, people tend to use only the upper part of the scale. Such tendencies can be observed on mean user rating distribution plot.


To correct for biases caused by varying mean ratings of different users and items(i.e. long or hard-to-watch movies can also be rated far lower than others) special factors are introduced in the form of user bias, item bias or baseline. [Section 5.2.1 Recommender Systems Handbook, Ricci]

```

In [14]: user_ratings_count_plot = sns.distplot(ratings_df.groupby('user_id')['rating'].count(),
        kde=False)
        user_ratings_count_plot.set(xlabel='Ratings count', ylabel='Frequency')

```

Out[14]: [Text(0, 0.5, 'Frequency'), Text(0.5, 0, 'Ratings count')]




```
In [15]: len(ratings_df.groupby('book_id')['rating'].count()[ratings_df.groupby('book_id')['rating'].count() < 10000])
```

```
Out[15]: 9958
```

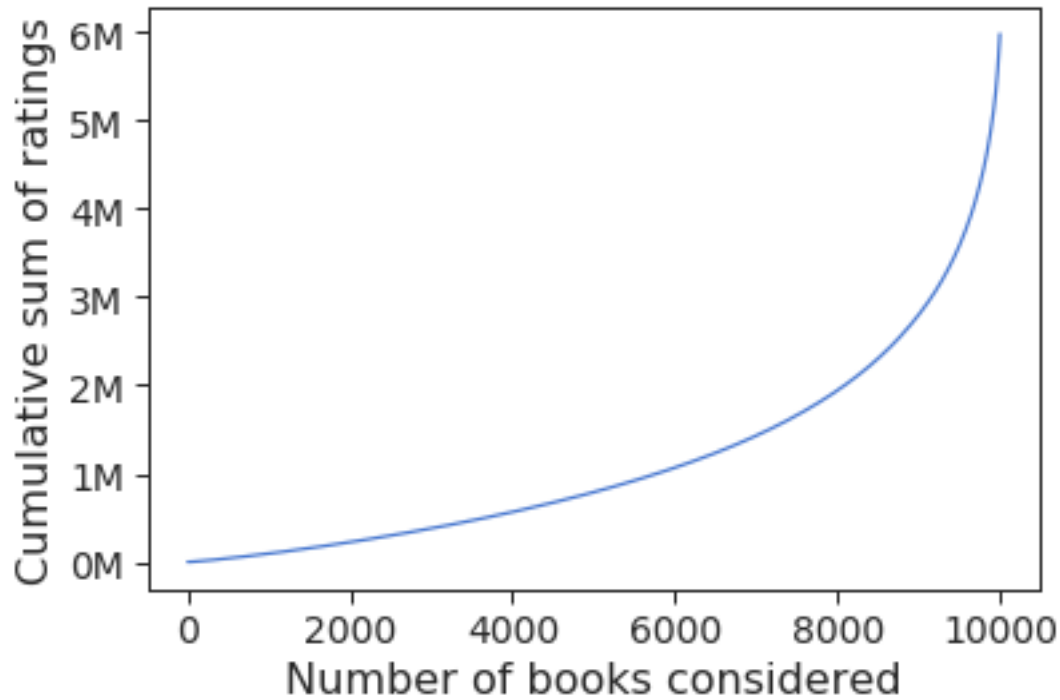
```
In [16]: book_ratings_count_plot = sns.distplot(ratings_df.groupby('book_id')['rating'].count(),
        kde=False)
        book_ratings_count_plot.set_yscale('log')
        book_ratings_count_plot.set(xlabel='Ratings count', ylabel='Frequency')
```

```
Out[16]: [Text(0, 0.5, 'Frequency'), Text(0.5, 0, 'Ratings count')]
```



```
In [17]: book_ratings_cum_count =
ratings_df.groupby('book_id')['rating'].count().sort_values().cumsum()

In [18]: book_ratings_cum_count_plot = sns.lineplot(y=book_ratings_cum_count.values, x=[x+1 for x
in range(0,10000)])
book_ratings_cum_count_plot.set(xlabel='Number of books considered', ylabel='Cumulative
sum of ratings')
book_ratings_cum_count_plot.yaxis.set_major_formatter(
matplotlib.ticker.FuncFormatter(lambda x, pos: '{:,.0f}'.format(x/(10**6)) + 'M'))
```



## 6 Train and test split

```
In [19]: from sklearn.model_selection import train_test_split
```

```
In [20]: train_df, test_df = train_test_split(ratings_df, test_size=0.1, random_state=44)
```

Some used methods do not generalize well for new(unseen) users and items, so we have to make sure that training test contains all users and items.

```
In [21]: set(train_df['user_id'].unique()) == set(ratings_df['user_id'].unique())
```

```
Out[21]: True
```

```
In [22]: set(train_df['book_id'].unique()) == set(ratings_df['book_id'].unique())
```

```
Out[22]: True
```

```
In [23]: train_df.groupby('user_id')['book_id'].count().describe()
```

```
Out[23]: count      53424.000000
         mean        100.681922
         std         23.671726
         min         17.000000
         25%         86.000000
         50%        100.000000
         75%        115.000000
         max        182.000000
         Name: book_id, dtype: float64
```

```
In [24]: train_df.groupby('book_id')['user_id'].count().describe()
```

```
Out[24]: count      10000.000000  
         mean         537.883100  
         std         1140.646885  
         min           8.000000  
         25%         140.000000  
         50%         223.000000  
         75%         454.250000  
         max        20508.000000  
         Name: user_id, dtype: float64
```